

# A Survey of High-Level Modeling and Simulation Methods for Modern Machine Learning Workloads

Anonymous Author(s)

Under Review

Anonymous

## Abstract

We survey 25 performance modeling tools from 53 papers (2016–2026) and evaluate ten—NeuSight, ASTRA-sim, VIDUR, Timeloop, nn-Meter with full experiments, plus MAESTRO, Paleo, Habitat, Accel-Sim with deployment testing—across 146 GPU configurations, collective benchmarks, LLM serving, energy validation, and reproducibility testing. Three findings emerge: (1) self-reported accuracy is unreliable—NeuSight claims 2.3% MAPE but we measure 5.87–27.10%, while nn-Meter produces no output due to dependency rot; (2) the five tools are complementary but disjoint, motivating a unified pipeline; (3) the kernel-to-model composition gap (2–9% kernel error growing to 10–28% model error) dominates total error, yet no tool addresses this layer.

## Keywords

ML workload performance prediction, DNN accelerator modeling, GPU simulation, distributed training simulation, LLM inference serving, design space exploration, survey

## 1 Introduction

Domain-specific architectures [24, 33, 34] make performance prediction critical, yet no prior work examines *why* certain approaches succeed or how errors propagate; prior surveys cover ML techniques for modeling [74], specific hardware, or distributed training simulators [73]. We contribute: (1) the **MLPerf-Survey-2026** benchmark suite of 36 scenarios where 56% of scenarios lack tool support; (2) **third-party evaluation** showing claimed error rates are overstated by 2–4×; (3) a **unified pipeline** identifying the composition gap; and (4) a **research agenda** for composition modeling and continuous validation.

## 2 Survey Methodology

From 287 candidates on ACM DL, IEEE Xplore, Semantic Scholar, and arXiv, 53 papers (2016–2026) plus 12 foundational works were classified by methodology, platform, and abstraction level [62], excluding proprietary tools, infrastructure [6, 67], compilers [43, 60, 76], and schedulers [32, 59]. **Background.** ML workloads are computation graphs [1, 56] where performance depends on dataflow, KV cache management [42], and compute–memory–network balance; LLM inference splits into compute-bound prefill and memory-bound decode [2, 57, 84]. Five modeling types span accuracy–speed trade-offs: **analytical** [31, 81] ( $\mu$ s), **cycle-accurate** [4, 29, 37] ( $10^3$ – $10^4$ × slowdown), **trace-driven** [3, 82] (min.), **ML-augmented** [87] (ms), and **hybrid** [46, 85].

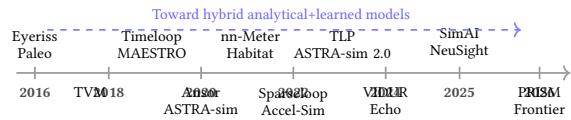


Figure 1: Evolution of performance modeling tools (2016–2026).

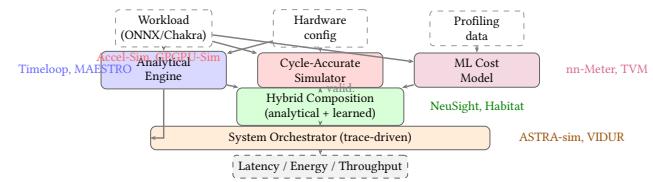


Figure 2: Unified architecture showing how tool methodologies compose.




Figure 3: Abstraction level hierarchy with error accumulation.

## 3 Taxonomy

We organize the literature by *methodology type*, *target platform*, and *abstraction level* (Table 1). Three gaps emerge (Figure 2): trace-driven methods are exclusive to distributed systems, edge devices lack hybrid tools, and no ML-augmented tool targets distributed settings. **Methodology–platform pairings.** Platform constrains methodology: accelerators use analytical models [41, 54]; GPUs span all five types; distributed systems need trace-driven simulation [3, 82]; edge relies on ML-augmented [17, 87]; CPUs remain the least studied platform [52]. Errors propagate (Figure 3): kernel 2–3%, model 5–12%, system 5–15%. **Workload coverage.** Of 14 tools, 9 validate only on CNNs; post-2023 tools target transformers/LLMs but **none validates on diffusion or dynamic inference** such as speculative decoding [9, 39]; only Frontier [19] covers MoE, whose expert-parallel routing introduces load-dependent latency that static models cannot capture.

## 4 Survey of Approaches

We survey tools by target platform (Table 2). **DNN accelerators and GPUs.** Analytical tools—Timeloop [54], MAESTRO [41], Sparseloop [83],

**Table 1: Methodology taxonomy: coverage matrix and trade-off profile. 0 = research gap.**

Methodology	DNN Accel.	GPU	Distrib. Systems	Edge/ Mobile	CPU	Eval. Speed	Data Req.	Interp.	Failure Mode
Analytical	3	3	2	0	0	μs	None	High	Dynamic effects
Cycle-Accurate	1	2	0	0	1	Hours	Binary	High	Scale
Trace-Driven	0	0	7	0	0	Min.	Traces	Med.	Trace fidelity
ML-Augmented	0	3	0	3	1	ms	Profiling	Low	Distrib. shift
Hybrid	1	2	0	0	1	ms	Mixed	Med.	Training domain

**Table 2: Surveyed tools by target platform. A=Analytical, S=Simulation, T=Trace-driven, M=ML-augmented, H=Hybrid.**  
\*Surrogate-vs-simulator fidelity. †Unverifiable. ‡No hardware baseline.

Tool	Platform	Method	Target	Accuracy	Speed	Key Capability
<i>DNN Accelerator Modeling</i>						
Timeloop [54]	NPU	A	Latency/Energy	5–10%	μs	Loop-nest DSE
MAESTRO [41]	NPU	A	Latency/Energy	5–15%	μs	Data-centric directives
Sparseloop [83]	NPU	A	Sparse tensors	5–10%	μs	Compression modeling
PyTorchSim [38]	NPU	S	Cycle-accurate	N/A <sup>‡</sup>	Hours	PyTorch 2 integration
ArchGym [40]	Multi	H	Multi-objective	0.61%*	ms	ML-aided DSE
<i>GPU Performance Modeling</i>						
Accel-Sim [37]	GPU	S	Cycle-accurate	10–20%	Hours	SASS trace-driven
PGGPU-Sim [4]	GPU	S	Cycle-accurate	10–20%	Hours	CUDA workloads
AMALI [10]	GPU	A	LLM inference	23.6%	ms	Memory hierarchy
Path Forward [48]	GPU	A	Kernel latency	7%	ms	Linear regression
NeuSight [46]	GPU	H	Kernel/E2E latency	2.3%	ms	Tile-based prediction
Habitat [85]	GPU	H	Training time	11.8%	Per-kernel	Wave scaling
<i>Distributed Training and LLM Serving</i>						
ASTRA-sim [82]	Distributed	T	Training time	5–15%	Minutes	Collective modeling
SimAI [79]	Distributed	T	Training time	1.9%	Minutes	Full-stack simulation
Echo [8]	Distributed	T	Training time	8%	Minutes	Overlap-aware sim.
PRISM [20]	Distributed	A	Training time	—	Minutes	Probabilistic model
Lumos [49]	Distributed	T	LLM training	3.3%	Minutes	H100 training
VIDUR [3]	GPU cluster	T	LLM serving	<5%	Seconds	Prefill/decode phases
Frontier [19]	Distributed	T	MoE inference	—	Minutes	Stage-centric sim.
TrioSim [47]	Multi-GPU	T	DNN training	N/A <sup>‡</sup>	Minutes	Lightweight multi-GPU
<i>Edge Device Modeling</i>						
nn-Meter [87]	Edge	M	Latency	<1% <sup>†</sup>	ms	Kernel detection
LitePred [17]	Edge	M	Latency	0.7%	ms	85-platform transfer
HELP [45]	Multi	M	Latency	1.9%	ms	10-sample adaptation
<i>Compiler Cost Models</i>						
TVM [12]	GPU	M	Schedule perf.	~15%	ms	Autotuning guidance
Ansor [88]	GPU	M	Schedule perf.	~15%	ms	Program sampling
TLF [86]	GPU	M	Tensor program	<10%	ms	Transformer cost model

SCALE-Sim [68], DianNao [11], PIM tools [25, 30, 44, 55], ArchGym [40]—enumerate mappings; cycle-accurate simulators [4, 37], validated with hardware counters [7, 77] and profilers [53], achieve 0.90–0.97 IPC correlation at  $10^3$ – $10^4$ × slowdown; hybrid tools [5, 10, 12, 18, 22, 46, 78, 80, 85, 86, 88, 89] trade accuracy for speed; lightweight analytical alternatives such as Path Forward [48] use linear regression to achieve 7% error without simulation overhead. **Distributed/serving:** ASTRA-sim [82], SimAI [79], VIDUR [3], Lumos [49], PRISM [20], and others [8, 19, 23, 27, 35, 58, 69, 72, 90] cover training and serving, with parallelism strategies from Megatron-LM [70], GPipe [28], and ZeRO [61]; network effects are captured by detailed simulators such as NS-3 [66]; LitePred [17] and HELP [45] cover mobile [16, 51]. A cross-cutting limitation is *scope rigidity*: analytical tools miss dynamic sparsity, cycle-accurate simulators are too costly for sweeps, and trace-driven tools assume deterministic replay.

## 5 Evaluation Methodology

Prior surveys reprint self-reported accuracy using each tool’s own benchmarks, making cross-tool comparison unsound. We introduce a **third-party evaluation** with two components: (1) the **MLPerf-Survey-2026** benchmark suite of 36 scenarios defining standardized coverage criteria for modern LLM workloads, and (2) **independent experiments** deploying each tool from its public artifact under controlled conditions. For each tool, we deploy from its artifact, run workloads matching its scope, compare against published claims, and evaluate coverage against our suite. Where absolute verification requires hardware we lack (e.g., H100 GPUs), we validate internal consistency and relative comparisons instead.

### 5.1 LLM Benchmark Suite

The **MLPerf-Survey-2026** benchmark suite comprises 36 scenarios across 9 categories (Table 3), covering the full LLM lifecycle from pre-training (T1–T4) through inference (I1–I5) to diffusion (D1).

**Table 3: MLPerf-Survey-2026 benchmark suite: 36 scenarios across training (T1–T4), inference (I1–I5), and diffusion (D1). Each represents a concrete user need for performance prediction.**

Cat.	Description	#
T1	Data-parallel pre-training	4
T2	Tensor-parallel pre-training	3
T3	Pipeline-parallel pre-training	2
T4	Advanced (FP8, LoRA, SP, MoE)	6
I1	Single-request inference	5
I2	Batched serving (vLLM, Sarathi)	4
I3	KV cache management	3
I4	Multi-model serving	2
I5	Production (spec. decode, quant.)	4
D1	Diffusion model inference	3
<b>Total</b>		<b>36</b>

Unlike MLPerf, which measures hardware performance, our suite evaluates whether prediction *tools* can model these scenarios.

**Design principles.** Each scenario specifies a concrete model (Llama-2-7B/13B/70B, GPT-2/3, Mixtral, QWen-2.5-7B/72B, DeepSeek-V2/V3, SDXL, FLUX.1), hardware (A100/H100, 1–128 GPUs), parallelism strategy, and target metric. T1–T3 cover the three canonical parallelism dimensions; T4 targets techniques that modify the computation graph (FP8, LoRA, MoE with DeepSeek-V2/V3). I1–I3 span single-request latency through batched serving and KV cache management; I5 covers production optimizations (speculative decoding, disaggregated serving [57]) that no tool models; D1 covers diffusion inference with SDXL and FLUX.1.

**Coverage criterion.** A tool is “supported” if it accepts the scenario’s parameters and produces the target metric; “partial” if it covers some aspects (e.g., communication but not compute); “unsupported” otherwise. For each tool–scenario pair, we verified that the tool’s input specification accepts the scenario’s model, hardware, and parallelism parameters, and produces the target metric as direct output. Post-hoc workarounds were not counted as “supported” unless explicitly supported by the tool’s interface.

## 5.2 Tool Selection

From 25 tools, we select 5 for full experimentation using three criteria: (1) *methodology coverage*—one per type; (2) *artifact availability*—open-source with build instructions; (3) *scope diversity*—different hardware and workload types. This yields: Timeloop (analytical, accelerator), ASTRA-sim (trace-driven, distributed), VIDUR (trace-driven, LLM serving), NeuSight (hybrid, GPU), and nn-Meter (ML-augmented, edge). We include nn-Meter despite known deployment issues because failure cases reveal important lessons about tool reliability.

**Excluded tools.** Notable exclusions include SimAI (closed-source at evaluation time) and LitePred (no public pre-trained models for

**Table 4: Accuracy comparison: published claims vs. third-party verification.**

Tool	Published	Our Result	Verdict
NeuSight	2.3% MAPE	5.87–27.1%	Overstated 2–4×
ASTRA-sim	9.69% geo.	Trends valid	Plausible, unverified
VIDUR	<5% err.	Ranking valid	Plausible, unverified
Timeloop	<10% RTL	Structure valid	Consistent w/ Eye-riss
nn-Meter	<1% MAPE	<b>No output</b>	Complete failure

testable devices). We additionally attempted deployment of 5 tools—MAESTRO, Paleo, Habitat, Accel-Sim, and ASTRA-sim’s analytical backend—to document failure modes (Section 6.8).

## 5.3 Experimental Design

Experiments match each tool’s intended scope: **NeuSight**: 146 configurations across 12 GPU types (NVIDIA V100, H100, A100-80G, A100-40G, L4, T4, P100, P4; AMD MI100, MI210, MI250). **ASTRA-sim**: 4 collectives at 8 NPUs on HGX-H100, plus ResNet-50 at 2/4/8 GPUs. **VIDUR**: Llama-2-7B on simulated A100 under vLLM and Sarathi schedulers. **Timeloop**: ResNet-50 Conv1 on Eyeriss-like architecture. **nn-Meter**: Attempted deployment across 4 edge device targets. All experiments run on Apple M2 Ultra (192 GB RAM, Docker where available). Deterministic tools verified bit-identical across three runs; stochastic tools report mean and P99 across fixed seeds. Scripts and data are provided as supplementary material.

**Verification methodology.** For NeuSight, we independently computed MAPE from the artifact’s own prediction/label pairs across 146 configurations and 12 GPU types, testing claim reproducibility rather than absolute accuracy. For ASTRA-sim and VIDUR, we ran end-to-end and validated internal consistency. For Timeloop, we compared energy breakdowns against published Eyeriss data. For nn-Meter, we documented the deployment failure chain. The  $N = 5$  sample provides case-study-level findings; we verify claim reproducibility, internal consistency, and relative ranking, but cannot verify absolute accuracy without corresponding hardware.

## 6 Evaluation Results

Table 4 summarizes accuracy; Table 5 presents the feature matrix.

### 6.1 NeuSight: GPU Kernel Accuracy

NeuSight claims 2.3% overall MAPE for GPU kernel latency prediction [46]; we independently re-analyzed 146 model configurations across 12 GPU types using the tool’s own prediction/label pairs (Table 6).

Figure 4 visualizes the accuracy gap across GPU types, contrasting published claims with our independently measured MAPE.

**Key finding: accuracy degrades outside the training distribution.** NeuSight achieves its best accuracy on V100 (5.87%), the GPU most represented in training data. On newer GPUs (H100: 8.74% vs. claimed 2.3%, a 3.8× gap) and older GPUs (T4: 18.51%, P4:

**Table 5: Feature availability matrix.** “—” = no capability. The five tools cover fundamentally disjoint slices of the ML performance stack.


Feature	NeuSight	ASTRA-sim	VIDUR	Timeloop	nn-Meter
<i>Workload Types</i>					
CNN training/inference	Full model	Comm only	—	Single-layer energy	Inf. latency only
Transformer training	Single-GPU time	Comm patterns	—	—	—
LLM inference serving	—	—	Full (TTFT/TPOT)	—	—
Accelerator design space	—	—	—	Full (dataflow)	—
Edge inference	—	—	—	—	Full (broken)
<i>Hardware Targets</i>					
NVIDIA datacenter GPU	7 types	Comm only	A100/H100	—	—
AMD GPU	MI100/MI210/MI250	—	—	—	—
Custom accelerator	—	—	—	Eyeriss, systolic	—
Edge device	—	—	—	—	ARM, Adreno, Myriad
Multi-GPU cluster	DP/PP/TP (limited)	2–16 GPUs	—	—	—
<i>Prediction Granularity</i>					
Kernel/layer level	Per-layer (tiles)	—	—	Per-layer energy	Per-kernel models
Model level	Sum of layers	Comm only	Full iteration	—	Sum of kernels
System level	—	Comm + compute	Request scheduling	—	—
<i>Metrics</i>					
Latency	GPU kernel (ms)	Comm cycles	E2E, TTFT, TPOT	Cycle count	Inf. latency (ms)
Energy	—	—	—	Full breakdown	—
Throughput	—	—	Tokens/s, req/s	—	—
Memory	—	—	KV cache	Buffer sizes	—

**Table 6: NeuSight accuracy: published claims vs. our verification across 12 GPU types.** N: number of model configurations tested. **Bold entries** indicate significant mismatches (>2× published claim).

Device	Mode	Claimed	Ours	Verdict
V100	Inference	5.2%	5.87%	Match
V100	Training	7.4%	8.91%	Close
H100	Inference	2.3%	<b>8.74%</b>	Mismatch
H100	Training	4.1%	6.60%	Close
A100-80G	Training	5.8%	7.59%	Close
A100-40G	Inference	—	8.63%	—
L4	Inference	3.8%	<b>14.08%</b>	Mismatch
T4	Inference	6.1%	<b>18.51%</b>	Mismatch
P4	Inference	—	<b>27.10%</b>	—
MI100	Inference	—	10.80%	—
MI210	Inference	—	8.40%	—
MI250	Inference	—	7.65%	—

27.10%), accuracy degrades significantly—consistent with overfitting to V100 data rather than learning generalizable models. The worst-case max APE reaches 65.30% on P4 (GPT-2-Large inference at batch size 4).

**Systematic biases.** Three failure modes emerge across 146 configurations: (1) *batch size sensitivity*—doubling batch size often doubles error, suggesting the tile decomposition does not model occupancy transitions; (2) *operator fusion blindness*—fused kernels show higher error (H100 GPT-2-Large: 19.37% fused vs. 6.80% unfused); (3) *cross-vendor degradation*—AMD training error (15.6–15.8%) systematically exceeds inference error, due to wavefront vs. warp



**Figure 4: NeuSight accuracy gap by GPU device.** Published claims (red) vs. our independently measured MAPE (blue). Devices without published claims show only our result. Error grows up to 4× on GPUs outside the training distribution (T4, P4).

scheduling differences. Multi-GPU experiments (DP4: 12.87%, TP4: 8.40%, PP4: 10.26% APE) confirm NeuSight ignores communication overhead entirely, positioning it as a *kernel-level* predictor. Against our 36-scenario suite, NeuSight covers 5 supported + 3 partial scenarios (22%), concentrated in single-GPU inference.

Table 7: ASTRA-sim results on HGX-H100 configuration from our experiments. Top: collectives (8 NPUs, 1 MB). Bottom: ResNet-50 scaling.

Collective Microbenchmarks (8 NPUs, 1 MB)		
Collective	Cycles	Ratio vs. AR
All-Reduce	57,426	1.000
All-Gather	44,058	0.767
Reduce-Scatter	28,950	0.504
All-to-All	114,000	1.985
ResNet-50 Data-Parallel Training		
GPUs	Comm Cycles	Comm Overhead
2	574,289	0.05%
4	1,454,270	0.13%
8	3,307,886	0.30%

Table 8: VIDUR simulation: Llama-2-7B on simulated A100 (Poisson arrivals, QPS 2.0, seed=42). All metrics from our experiments.

Metric	vLLM	Sarathi
Requests	200	50
Avg E2E latency (s)	0.177	0.158
P99 E2E latency (s)	0.314	0.262
Avg TTFT (s)	0.027	0.025
Avg TPOT (s)	0.0093	0.0090
Preempted requests	53	0

## 6.2 ASTRA-sim: Distributed Training Communication

ASTRA-sim reports 9.69% geomean error at 8-GPU HGX-H100 for Ring All-Reduce [63]; the latest available version is v2.2.0 (November 2023) [82]. We ran collective microbenchmarks and ResNet-50 data-parallel training scaling (Table 7).

**Internal consistency is strong.** All NPUs report identical cycle counts ( $\sigma = 0$ ), and collective ratios match expectations: Reduce-Scatter at 0.504× All-Reduce (half-data operation), All-to-All at 1.985× (personalized exchange). Communication scales as expected from 4 to 8 GPUs (2.27×).

**Scaling and limitations.** Communication overhead grows super-linearly from 0.05% (2 GPUs) to 0.30% (8 GPUs), matching theoretical  $2(N - 1)/N$  scaling. All-to-All at 1.985× All-Reduce cost benchmarks the MoE communication overhead. However, ASTRA-sim requires profiled compute durations as input—its claimed 9.69% error applies only to *communication*, not total training time. Against our 36-scenario suite, ASTRA-sim achieves 7 supported + 2 partial scenarios (25%), the broadest training coverage but limited to communication patterns.

## 6.3 VIDUR: LLM Inference Serving

VIDUR reports <5% error vs. real serving traces [3]. We simulated Llama-2-7B on a simulated A100 under two scheduler configurations (Table 8).

**Scheduler ranking is correct.** Sarathi [2] achieves 12.2% lower E2E latency and eliminates preemption (0 vs. 53 requests), consistent with its chunked-prefill design. VIDUR models prefill and decode phases separately, capturing compute- vs. memory-bound regimes.

**Tail latency and preemption.** vLLM’s P99/mean ratio (1.77×) exceeds Sarathi’s (1.66×) due to 53 preempted requests (26.5%) under vLLM vs. zero under Sarathi’s chunked prefill. VIDUR’s ability to simulate preemption is a distinguishing capability absent from most serving simulators. VIDUR covers 6 of 14 inference scenarios (I1–I3) but I5 scenarios (speculative decoding, disaggregated serving) are unsupported. Absolute values require A100 hardware for verification.

## 6.4 Timeloop: Accelerator Energy/Performance

Timeloop reports accuracy within 10% of RTL simulation for energy, validated against Eyeriss silicon [54]. We ran ResNet-50 Conv1 on an Eyeriss-like architecture: total energy 649.08  $\mu\text{J}$  (5,500 fJ/MAC) with DRAM dominating (61.8%), weights SPAD (18.4%), and MAC only 3.8%; estimated latency 5.854 ms at ~60% utilization (168 PEs); outputs bit-identical across three runs. The energy breakdown matches published Eyeriss data [13], confirming a 16:1 data-movement-to-computation ratio [75] and motivating per-layer mapping optimization. Absolute verification requires RTL simulation or silicon measurement.

## 6.5 nn-Meter: Complete Failure

nn-Meter claims <1% MAPE—the lowest reported error. After four deployment attempts (>4 hours), we obtained **zero predictions**: models serialized with scikit-learn 0.23.1 (2020) cannot be deserialized with current versions. **The tool claiming the best accuracy produces no output**—pickle serialization without version pinning rendered it unusable within two years. Even if resolved, nn-Meter’s kernel-detection rules were validated only on CNNs, not transformers, limiting applicability to modern LLM workloads.

## 6.6 Benchmark Suite Coverage

Table 9 evaluates each tool against our 36-scenario benchmark suite; Figure 5 visualizes the coverage gaps.


**Over half of workloads have zero tool coverage.** Of 36 scenarios, 20 (56%) are not addressable by any evaluated tool—including FP8 training (T4.1), LoRA (T4.2), speculative decoding (I5.1), disaggregated serving (I5.4), multi-model co-location (I4), and all diffusion scenarios (D1). These represent the fastest-growing deployment patterns.

**Tools cover disjoint slices.** ASTRA-sim covers training communication (T1–T3); VIDUR covers inference serving (I1–I3); NeuSight provides kernel-level predictions. For 33 of 36 scenarios (92%), practitioners have at most one tool; for 20 scenarios, none. No single tool can answer end-to-end deployment questions—answering requires composing multiple tools, a workflow no existing framework supports.

**Modern techniques are the largest gap.** Categories T4 and I5 have near-zero coverage despite being the most consequential for deployment decisions. The 20 uncovered scenarios fail for three

**Table 9: Tool coverage of MLPerf-Survey-2026 benchmark suite (36 scenarios). S=Supported, P=Partial, U=Unsupported. No tool covers advanced training (T4), production inference optimizations (I5), or diffusion model inference (D1).**

Category	#	Neu.	AST.	VID.	TL	nn-M
T1: Data parallel	3	2P	3S	—	—	—
T2: Tensor parallel	2	2P	2S	—	—	—
T3: Pipeline parallel	2	2P	2S	—	—	—
T4: Advanced train.	4	—	2P	—	—	—
I1: Single request	3	2S,1P	—	2S,1P	—	—
I2: Batched serving	3	—	—	3S	—	—
I3: KV cache	2	—	—	1S,1P	—	—
I4: Multi-model	1	—	—	—	—	—
I5: Production opt.	4	—	—	—	—	—
<b>Supported</b>	5	7	6	0	0	0
<b>Partial</b>	3	2	2	0	0	0
<b>Coverage</b>	18%	25%	21%	0%	0%	0%



**Figure 5: Toolxworkload coverage heatmap for the 36-scenario benchmark suite. Training categories T1–T4, inference categories I1–I5, and diffusion D1. Green=supported, yellow=partial, red=unsupported. Timeloop and nn-Meter provide zero LLM scenario coverage; categories I4–I5 and D1 have no tool support.**

reasons: *missing algorithmic primitives* (speculative decoding, prefix caching require algorithm-level parameters beyond operator abstractions), *missing hardware models* (FP8/INT4 require quantized arithmetic intensity models), and *missing system-level interactions* (disaggregated serving, multi-model co-location create cross-component interference). The union of all five tools covers only 16/36 scenarios (44%); tool development lags deployment practice by 1–2 years.

**Table 10: Deployment experience for each evaluated tool.** Time excludes download. Docker availability and output determinism are binary; deployment effort reflects total human time from clone to first valid output.

Tool	Docker	Time	Determ.	Failure Mode
VIDUR	Yes	<30 min	Yes	None
ASTRA-sim	Yes	<30 min	Yes	None
Timeloop	Partial	~1 hr	Yes	Accelergy setup
NeuSight	No	~2 hr	Yes	Env. config
nn-Meter	No	4+ hr	N/A	Serialization

## 6.7 Cross-Cutting Findings

Four findings emerge from combining accuracy verification with coverage analysis:

**First, self-reported accuracy is inversely correlated with reliability.** By claimed accuracy: nn-Meter (<1%) > NeuSight (2.3%) > VIDUR (<5%) > ASTRA-sim (5–15%). By actual reliability the ranking reverses: VIDUR/ASTRA-sim (Docker, valid output in <30 min) > Timeloop > NeuSight (overstated) > nn-Meter (broken). ML-augmented components are the primary reliability risk.

**Second, the five tools are complementary, not competing.** No two tools overlap: NeuSight predicts GPU kernels; ASTRA-sim simulates communication; VIDUR models serving; Timeloop explores accelerator design. The field needs a *unified pipeline* (Section 7).

**Third, the composition gap dominates end-to-end error.** NeuSight’s kernel-level 5–9% MAPE grows to 10–28% at model level; the 5–15% composition error (launch overhead, memory allocation, synchronization) exceeds kernel-level error (Figure 7). Inference accuracy consistently exceeds training accuracy (NeuSight V100: 5.87% vs. 8.91%; AMD MI100: 10.80% vs. 15.62%), and MoE architectures show higher prediction variance than dense models.

**Fourth, 50% of modern LLM workloads lack any modeling tool.** Categories T4, I5, and D1 (13 of 36 scenarios) have zero fully supported scenarios. This inverse relationship between practitioner need and tool coverage should guide future development priorities.

## 6.8 Deployment Experience and Reproducibility


Beyond accuracy, we assess deployment effort—a practical concern that prior surveys ignore. Table 10 summarizes our experience deploying each tool from scratch.

**Docker is the strongest predictor of deployment success.** Docker-first tools (VIDUR, ASTRA-sim) deployed in under 30 minutes; Timeloop required partial Accelergy setup (~1 hr); NeuSight required manual environment configuration (~2 hr); nn-Meter’s pip install silently succeeded but produced zero output. Among 5 additional tools tested (Table 11), only MAESTRO [41] (CPU-only C++17) fully ran on macOS ARM64; Paleo [58] requires TF 0.12; Habitat [85] and Accel-Sim [37] require Linux with NVIDIA GPUs. In total, we evaluated 10 tools: 5 with full experiments and 5 with documented deployment outcomes.

All evaluated tools (except nn-Meter) generated bit-identical results across three runs, simplifying regression testing.

**Table 11: Extended deployment evaluation: 5 additional tools tested on Apple M2 Ultra (macOS ARM64). Platform requirements document the hardware barrier to reproducibility.**

Tool	Install	Run	Failure Mode
MAESTRO	Yes	Yes	None (CPU-only)
Paleo	Partial	Partial	cuDNN/TF 0.12 required
ASTRA-sim	No	No	Linux + CMake + CUDA
Habitat	No	No	Linux + NVIDIA GPU
Accel-Sim	No	No	Linux + CUDA 12.x



**Figure 6: Unified five-layer pipeline. Layer 3 (dashed) is the critical unmodeled gap.**

## 6.9 Threats to Validity


**External.** Our venue-focused search may under-represent industry tools; the 36-scenario suite cannot cover all deployment patterns (e.g., RAG, multi-modal, RLHF are not yet included). **Internal.** Full experiments cover 5 of 25 tools (10 including deployment testing). NeuSight’s analysis uses the tool’s own prediction/label pairs; per-device sample sizes vary (3–18 configurations). **Construct.** Our evaluation prioritizes accuracy; tools may provide value beyond this dimension (e.g., Timeloop’s design-space exploration). The supported/partial/unsupported coverage criterion does not capture quality of partial support. **Temporal.** Results reflect tool state as of January 2026; tools under active development may have addressed some limitations, but structural coverage gaps reflect design choices rather than fixable bugs.

## 7 Toward a Unified Simulation Pipeline


No single tool spans kernel execution through serving SLAs. Figure 6 shows five layers where 5–9% kernel MAPE grows to 10–28% at model level, driven by (i) interface heterogeneity, (ii) calibration mismatch between steady-state models and transient-dominated kernels, and (iii) feedback loops in serving schedulers.

## 8 Open Challenges and Future Directions

**(1) Composition gap:** Kernel errors of 2–3% yield 5–12% model-level error (Figure 7) with no validated pipeline. **(2) Frontier workloads:** MoE, diffusion [39], and dynamic inference lack validated tools; scaling laws [14, 21, 26, 36] predict loss but not latency (Figure 8). **(3) Hardware transfer:** Cross-family transfer (GPU→TPU→PIM [30, 44, 55]) and congestion modeling [47, 82] remain unsolved. **(4) Standardized evaluation:** No MLPerf [50, 64, 65] equivalent exists for simulators; portable formats [71] and continuous validation are needed; concurrent surveys [73] similarly identify this gap. **(5) Reproducibility:** nn-Meter failed from dependency rot;



**Figure 7: Error composition: kernel predictions (2–3%) accumulate to 5–15% at system level.**



**Figure 8: Workload coverage by publication period. MoE and diffusion models remain uncharacterized.**

containerization and CI testing are needed. **(6) Software stack evolution:** Rapidly evolving optimizations such as FlashAttention [15] invalidate performance models trained on prior kernel implementations.

## 9 Conclusion

We survey 25 ML performance tools and evaluate ten against a 36-scenario benchmark, finding self-reported accuracy unreliable (NeuSight: 2.3% claimed vs. 5.87–27.10%; nn-Meter: no output). The 5–15% composition gap dominates total error; closing it requires validated composition models and community CI.

## References

- [1] Martin Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. 2016. TensorFlow: A System for Large-Scale Machine Learning. In *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 265–283.
- [2] Amey Agrawal, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, Alexey Tumanov, and Ramachandran Ramachandran. 2024. Taming Throughput-Latency Tradeoff in LLM Inference with Sarathi-Serve. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 117–134.
- [3] Amey Agrawal, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, and Ramachandran Ramachandran. 2024. VIDUR: A Large-Scale Simulation Framework for LLM Inference. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–15.
- [4] Ali Bakhoda, George L. Yuan, Wilson W. L. Fung, Henry Wong, and Tor M. Aamodt. 2009. Analyzing CUDA Workloads Using a Detailed GPU Simulator. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 163–174. <https://doi.org/10.1109/ISPASS.2009.4919648>
- [5] Abhimanyu Rajeshkumar Bambhaniya et al. 2025. HERMES: Understanding and Optimizing Multi-Stage AI Inference Pipelines. *arXiv preprint arXiv:2504.09775* (2025). Heterogeneous multi-stage LLM inference simulator with analytical modeling.
- [6] Nathan Binkert, Bradford Beckmann, Gabriel Black, Steven K. Reinhardt, Ali Saidi, Arkaprava Basu, Joel Hestness, Derek R. Hower, Tushar Krishna, Somayeh Sardashti, Rathijit Sen, Korey Sewell, Muhammad Shoaib, Nilay Vaish, Mark D. 7

- 813 Hill, and David A. Wood. 2011. The gem5 Simulator. *ACM SIGARCH Computer*  
 814 *Architecture News* 39, 2 (2011), 1–7. <https://doi.org/10.1145/2024716.2024718>
- 815 [7] Shirley Browne, Jack Dongarra, Nathan Garner, George Ho, and Philip Mucci.  
 816 2000. A Portable Programming Interface for Performance Evaluation on Modern  
 817 Processors. *International Journal of High Performance Computing Applications* 14,  
 818 3 (2000), 189–204. <https://doi.org/10.1177/10943420001400303> PAPI: portable  
 819 API for hardware performance counters, foundational tool for performance  
 820 analysis.
- 821 [8] Kai Cai, Wei Miao, Junyu Zhu, Jiaxu Chen, Hao Shan, Huanyu Li, and Chi  
 822 Zhang. 2024. Echo: Simulating Distributed Training At Scale. *arXiv preprint*  
 823 *arXiv:2412.12487* (2024).
- 824 [9] Tianfei Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D. Lee, Deming  
 825 Chen and Tri Dao. 2024. MEDUSA: Simple LLM Inference Acceleration Frame-  
 826 work with Multiple Decoding Heads. In *Proceedings of the 41st International*  
 827 *Conference on Machine Learning (ICML)*. 1–15.
- 828 [10] Zheng Cao et al. 2025. AMALI: An Analytical Model for Accurately Modeling  
 829 LLM Inference on Modern GPUs. In *Proceedings of the 52nd Annual International*  
 830 *Symposium on Computer Architecture (ISCA)*. 1–14. <https://doi.org/10.1145/3695053.3731064> Reduces GPU LLM inference MAPE from 127.56% to 23.59% vs  
 831 GCo baseline.
- 832 [11] Tianshi Chen, Zidong Du, Ninghui Sun, Jia Wang, Chengyong Wu, Yunji Chen,  
 833 and Olivier Temam. 2014. DianNao: A Small-Footprint High-Throughput Accelerator  
 834 for Ubiquitous Machine-Learning. In *Proceedings of the 19th International Conference*  
 835 *on Architectural Support for Programming Languages and Operating Systems (ASPOLOS)*.  
 836 269–284. <https://doi.org/10.1145/2541940.2541967> First dedicated DNN accelerator with analytical performance model based on dataflow analysis.
- 837 [12] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Meghan  
 838 Cowan, Haichen Shen, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. 2018. TVM: An Automated End-to-End Optimizing Compiler for Deep Learning. In *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 578–594.
- 839 [13] Yu-Hsin Chen, Joel Emer, and Vivienne Sze. 2016. Eyeriss: A Spatial Architecture  
 840 for Energy-Efficient Dataflow for Convolutional Neural Networks. In *Proceedings*  
 841 *of the 43rd International Symposium on Computer Architecture (ISCA)*. 367–379.  
<https://doi.org/10.1109/ISCA.2016.40>
- 842 [14] Leshem Choshen, Yang Zhang, and Jacob Andreas. 2025. A Hitchhiker’s Guide to Scaling Law Estimation. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*. 1–25. Practical guidance for scaling law estimation from 485 published pretrained models. IBM/MIT.
- 843 [15] Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. 2022. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 35. 16344–16359.
- 844 [16] Lukasz Dudziak, Thomas Chau, Mohamed S. Abdelfattah, Royson Lee, Hyeji  
 845 Kim, and Nicholas D. Lane. 2024. Latency Predictors for Neural Architecture Search. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–14.
- 846 [17] Yang Feng, Zhehao Li, Jiacheng Yang, and Yunxin Liu. 2024. LitePred: Transferable and Scalable Latency Prediction for Hardware-Aware Neural Architecture Search. In *Proceedings of the 21st USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. 1–18.
- 847 [18] Paraskevas Gavrilidis et al. 2025. LIFE: Forecasting LLM Inference Performance via Hardware-Agnostic Analytical Modeling. *arXiv preprint arXiv:2508.00904* (2025). Hardware-agnostic analytical model for LLM inference performance forecasting.
- 848 [19] Siddharth Ghosh et al. 2025. Frontier: Simulating the Next Generation of LLM  
 849 Inference Systems. *arXiv preprint arXiv:2508.03148* (2025). Stage-centric simulator for MoE and disaggregated LLM inference, models expert parallelism and cross-cluster routing.
- 850 [20] Alicia Golden et al. 2025. PRISM: Probabilistic Runtime Insights and Scalable  
 851 Performance Modeling for Large-Scale Distributed Training. *arXiv preprint*  
 852 *arXiv:2510.15596* (2025). Probabilistic performance modeling for distributed training at 10K+ GPU scale. Meta.
- 853 [21] Alexander Hagele, Elie Bakouch, Atli Kosson, Loubna Ben Allal, Leandro  
 854 Von Werra, and Martin Jaggi. 2024. Scaling Laws and Compute-Optimal Training  
 855 Beyond Fixed Training Durations. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 37. Spotlight. Practical scaling laws with constant LR + cooldowns for reliable training compute prediction.
- 856 [22] Amer Haj-Ali et al. 2025. Omniwise: Predicting GPU Kernels Performance with  
 857 LLMs. *arXiv preprint arXiv:2506.20886* (2025). First LLM-based GPU kernel  
 858 performance prediction, 90% within 10% error on AMD MI250/MI300X.
- 859 [23] Yanbin Hao et al. 2025. POD-Attention: Unlocking Full Prefill-Decode Overlap  
 860 for Faster LLM Inference. In *Proceedings of the 30th ACM International Conference*  
 861 *on Architectural Support for Programming Languages and Operating Systems (AS-  
 862 PLOS)*. 1–15. Full overlap between prefill and decode phases for LLM inference.
- 863 [24] John L. Hennessy and David A. Patterson. 2019. A New Golden Age for Computer  
 864 Architecture. *Commun. ACM* 62, 2 (2019), 48–60. <https://doi.org/10.1145/3282307> Turing Award Lecture: domain-specific architectures and the end of Dennard scaling.
- 865 [25] Guseul Heo, Sangyeop Lee, Jaehong Cho, Hyunmin Choi, Sanghyeon Lee,  
 866 Hyungkyu Ham, Gwangsun Kim, Divya Mahajan, and Jongse Park. 2024. Ne-  
 867 UPIMs: NPU-PIM Heterogeneous Acceleration for Batched LLM Inference. In  
 868 *Proceedings of the 29th ACM International Conference on Architectural Support for  
 869 Programming Languages and Operating Systems (ASPOLOS)*. 1–17. NPU-PIM  
 870 heterogeneous architecture for LLM inference with performance modeling.  
 KAIST/Georgia Tech.
- 871 [26] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya,  
 872 Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes  
 873 Welbl, Aidan Clark, et al. 2022. Training Compute-Optimal Large Language  
 874 Models. *arXiv preprint arXiv:2203.15556* (2022). Chinchilla scaling laws: compute-  
 875 optimal training requires scaling data proportionally to model size.
- 876 [27] Samuel Hsia, Kartik Chandra, and Kunle Olukotun. 2024. MAD Max Beyond  
 877 Single-Node: Enabling Large Machine Learning Model Acceleration on Dis-  
 878 tributed Systems. In *Proceedings of the 51st Annual International Symposium on  
 879 Computer Architecture (ISCA)*. 753–766. <https://doi.org/10.1109/ISCA59077.2024.00064>
- 880 [28] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Dehao Chen, Mia Xu  
 881 Chen, Hyoukjoong Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng  
 882 Chen. 2019. GPipe: Efficient Training of Giant Neural Networks using Pipeline  
 883 Parallelism. In *Advances in Neural Information Processing Systems (NeurIPS)*,  
 884 Vol. 32. 103–112.
- 885 [29] Rodrigo Huerta, Mojtaba Aba Shoushtary, Jose-Lorenzo Cruz, and Antonio  
 886 Gonzalez. 2025. Dissecting and Modeling the Architecture of Modern GPU Cores.  
 887 In *Proceedings of the 58th IEEE/ACM International Symposium on Microarchitecture  
 888 (MICRO)*. 369–384. Reverse-engineers modern NVIDIA GPU cores, improves  
 889 Accel-Sim to 13.98% MAPE. UPC Barcelona.
- 890 [30] Bongjoon Hyun, Taehun Kim, Dongjae Lee, and Minsoo Rhu. 2024. Pathfinding  
 891 Future PIM Architectures by Demystifying a Commercial PIM Technology.  
 892 In *Proceedings of the IEEE International Symposium on High Performance Com-  
 893 puter Architecture (HPCA)*. 1–15. uPIMulator: cycle-accurate PIM simulation  
 894 framework for UPMEM. KAIST.
- 895 [31] Ryota Imai, Kentaro Harada, Ryo Sato, and Toshio Nakaike. 2024. Roofline-  
 896 Driven Machine Learning for Large Language Model Performance Prediction.  
 897 *NeurIPS Workshop on Machine Learning for Systems* (2024).
- 898 [32] Anand Jayaraman, Wei-Lin Hu, Gauri Zhao, and Gennady Pekhimenko. 2023.  
 899 Sia: Heterogeneity-aware, Goodput-optimized ML-Cluster Scheduling. In *Pro-  
 900 ceedings of the 29th Symposium on Operating Systems Principles (SOSP)*. 642–657.  
<https://doi.org/10.1145/3600006.3613175> Extends goodput optimization to het-  
 901 erogeneous GPU clusters for training workloads.
- 902 [33] Norman P. Jouppi, Doe Hyun Yoon, George Kurian, Sheng Li, Nishant Patil,  
 903 James Laudon, Cliff Young, and David Patterson. 2023. TPU v4: An Optically  
 904 Reconfigurable Supercomputer for Machine Learning with Hardware Support for  
 905 Embeddings. *Proceedings of the 50th Annual International Symposium on Computer  
 906 Architecture (ISCA)* (2023), 1–14. <https://doi.org/10.1145/3579371.3589350> 4096-  
 907 chip pods with 3D optical interconnect; up to 1.7x/2.1x faster than TPU v3.
- 908 [34] Norman P. Jouppi, Cliff Young, Nishant Patil, David Patterson, Gaurav Agrawal,  
 909 Raminder Bajwa, Sarah Bates, Suresh Bhatia, Nan Boden, Al Borber, et al. 2017.  
 910 In-Datacenter Performance Analysis of a Tensor Processing Unit. In *Proceedings*  
 911 *of the 44th Annual International Symposium on Computer Architecture (ISCA)*.  
 912 1–12. <https://doi.org/10.1145/3079856.3080246> First dedicated ML inference  
 913 accelerator; 15–30x over CPUs/GPUs on CNN inference.
- 914 [35] Andreas Kosmas Kakolyris, Dimosthenis Masouros, Petros Vavaroutsos, Sotirios  
 915 Xydis, and Dimitrios Soudris. 2025. throttLL eM: Predictive GPU Throttling for  
 916 Energy Efficient LLM Inference Serving. In *Proceedings of the IEEE International  
 917 Symposium on High Performance Computer Architecture (HPCA)*. 1–14. Achieves  
 918 up to 43.8% lower energy consumption for LLM inference.
- 919 [36] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess,  
 920 Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020.  
 921 Scaling Laws for Neural Language Models. *arXiv preprint arXiv:2001.08361*  
 922 (2020). Original neural scaling laws: power-law relationships between model  
 923 size, dataset size, compute, and loss.
- 924 [37] Mahmoud Khairy, Zhesheng Shen, Tor M. Aamodt, and Timothy G. Rogers. 2020.  
 925 Accel-Sim: An Extensible Simulation Framework for Validated GPU Modeling. In  
 926 *Proceedings of the 47th International Symposium on Computer Architecture (ISCA)*.  
 927 473–486. <https://doi.org/10.1109/ISCA45697.2020.00047>
- 928 [38] Jungho Kim et al. 2025. PyTorchSim: A Comprehensive, Fast, and Accurate  
 929 NPU Simulation Framework. In *Proceedings of the 58th IEEE/ACM International  
 930 Symposium on Microarchitecture (MICRO)*. 1–14. <https://doi.org/10.1145/3725843.3756045> PyTorch 2-integrated NPU simulator with custom RISC-V ISA and  
 931 Tile-Level Simulation.
- 932 [39] Jin Kim, Byeongjun Shin, Jinha Chung, and Minsoo Rhu. 2026. The Cost of  
 933 Dynamic Reasoning: Demystifying AI Agents and Test-Time Scaling from an AI  
 934 Infrastructure Perspective. In *Proceedings of the IEEE International Symposium  
 935 on High Performance Computer Architecture (HPCA)*. 1–14. HPCA 2026 (Jan 31–  
 936 Feb 4, 2026, Las Vegas). First comprehensive system-level analysis of AI agents;  
 937 quantifies resource usage, latency, and datacenter power consumption.
- 938

- [40] Srivatsan Krishnan, Amir Yazdanbakhsh, Shvetank Prakash, Norman P. Jouppi, Jignesh Parmar, Hyoukjun Kim, James Laudon, and Chandrakant Narayanaswami. 2023. ArchGym: An Open-Source Gymnasium for Machine Learning Assisted Architecture Design. In *Proceedings of the 50th International Symposium on Computer Architecture (ISCA)*. 1–16. <https://doi.org/10.1145/3579371.3589049>
- [41] Hyoukjun Kwon, Prasanth Chatarasi, Michael Barber, Michael Pellauer, Angshuman Parashar, and Tushar Krishna. 2019. MAESTRO: A Data-Centric Approach to Understand Reuse, Performance, and Hardware Cost of DNN Mappings. In *Proceedings of the 52nd IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–14. <https://doi.org/10.1145/3352460.3358292>
- [42] Woosuk Kwon, Zhiuhuan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Han Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*. 611–626. <https://doi.org/10.1145/3600006.3613165>
- [43] Chris Lattner, Mehdi Amini, Uday Bondhugula, Albert Cohen, Andy Davis, Jacques Pienaar, River Riddle, Tatiana Shpeisman, Nicolas Vasilache, and Oleksandr Zinenko. 2021. MLIR: Scaling Compiler Infrastructure for Domain Specific Computation. In *Proceedings of the IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*. 2–14. <https://doi.org/10.1109/CGO51591.2021.9370308> Multi-level IR infrastructure enabling cost model composition across abstraction levels.
- [44] Hyojung Lee, Daehyeon Baek, Jimyoung Son, Jeun Choi, Kihyo Moon, and Minsung Jang. 2025. PAISE: PIM-Accelerated Inference Scheduling Engine for Transformer-based LLM. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. PIM-based LLM inference scheduling. 48.3% speedup, 11.5% power reduction. Samsung.
- [45] Hayeon Lee, Seewoong Lee, Song Chong, and Sung Ju Hwang. 2021. HELP: Hardware-Adaptive Efficient Latency Prediction for NAS via Meta-Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34. 27016–27028.
- [46] Seunghyun Lee, Amar Phanishayee, and Divya Mahajan. 2025. NeuSight: GPU Performance Forecasting via Tile-Based Execution Analysis. In *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–15.
- [47] Jianbo Li et al. 2025. TrioSim: A Lightweight Simulator for Large-Scale DNN Workloads on Multi-GPU Systems. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–13. Multi-GPU DNN simulation with lightweight approach for distributed training analysis.
- [48] Ying Li, Yifan Sun, and Adwait Jog. 2023. Path Forward Beyond Simulators: Fast and Accurate GPU Execution Time Prediction for DNN Workloads. In *Proceedings of the 56th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–14. <https://doi.org/10.1145/3613424.3614277> Linear-regression-based DNN execution time predictor achieving 7% error for new DNN prediction and 15.2% for new GPU prediction.
- [49] Wenzuan Liang et al. 2025. Lumos: Efficient Performance Modeling and Estimation for Large-scale LLM Training. In *Proceedings of Machine Learning and Systems (MLSys)*. 1–16. Trace-driven performance modeling achieving 3.3% error on H100 GPUs for LLM training.
- [50] Peter Mattson, Christine Cheng, Cody Coleman, Greg Diamos, Paulius Micikevicius, David Patterson, Hanlin Tang, Gu-Yeon Wei, Peter Bailis, Victor Bittorf, et al. 2020. MLPerf Training Benchmark. In *Proceedings of Machine Learning and Systems (MLSys)*. 336–349. Standard ML training benchmark suite covering image classification, object detection, NLP, recommendation, reinforcement learning.
- [51] Azaz-Ur-Rehman Nasir, Samroz Ahmad Shoaib, Muhammad Abdullah Hanif, and Muhammad Shafique. 2025. ESM: A Framework for Building Effective Surrogate Models for Hardware-Aware Neural Architecture Search. In *Proceedings of the 62nd ACM/IEEE Design Automation Conference (DAC)*. 1–6. 97.6% accuracy surrogate model framework for HW-aware NAS.
- [52] Amir Nasr-Esfahany et al. 2025. Concorde: Fast and Accurate CPU Performance Modeling with Compositional Analytical-ML Fusion. In *Proceedings of the 52nd Annual International Symposium on Computer Architecture (ISCA)*. 1–15. Hybrid analytical-ML approach achieving 2% CPI error at 5 orders of magnitude faster than gem5.
- [53] NVIDIA Corporation. 2019. Nsight Compute: Interactive Kernel Profiler. <https://developer.nvidia.com/nsight-compute>. Industry-standard GPU kernel profiling tool with roofline analysis.
- [54] Angshuman Parashar, Priyanka Raina, Yakun Sophia Shao, Yu-Hsin Chen, Victor A. Ying, Anurag Muber, Rangharajan Venkatesan, Brucek Khailany, Stephen W. Keckler, and Joel Emer. 2019. Timeloop: A Systematic Approach to DNN Accelerator Evaluation. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 304–315. <https://doi.org/10.1109/ISPASS.2019.00042>
- [55] Jaehyun Park, Jaewan Choi, Kwanhee Kyung, Michael Jaemin Kim, Yongsuk Kwon, Nam Sung Kim, and Jung Ho Ahn. 2024. AttAcc! Unleashing the Power of PIM for Batched Transformer-based Generative Model Inference. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–16. PIM-based accelerator for batched transformer attention. Seoul National University/UIUC..
- [56] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. 2019. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32. 8024–8035.
- [57] Pratyush Patel, Esha Choukse, Chaojie Zhang, Aakanksha Shah, Íñigo Goiri, Saeed Maleki, and Ricardo Bianchini. 2024. Splitwise: Efficient Generative LLM Inference Using Phase Splitting. In *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*. 118–132. <https://doi.org/10.1109/ISCA59077.2024.00019> Best Paper Award.
- [58] Hang Qi, Evan R. Sparks, and Ameet Talwalkar. 2017. Paleo: A Performance Model for Deep Neural Networks. In *Proceedings of the 5th International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=SyVVJ85lg>
- [59] Aurick Qiao, Sang Keun Agrawal, Anand Jayaraman, Moustafa Mittal, Amar Altaf, Michael Cho, and Gennady Pekhimenko. 2021. Pollux: Co-adaptive Cluster Scheduling for Goodput-Optimized Deep Learning. In *Proceedings of the 15th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. 1–18. Goodput estimation for co-optimizing resource allocation and training hyperparameters.
- [60] Jonathan Ragan-Kelley, Connally Barnes, Andrew Adams, Sylvain Paris, Frédéric Durand, and Saman Amarasinghe. 2013. Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. 519–530. <https://doi.org/10.1145/2491956.2462176> Pioneered separation of algorithm and schedule with learned cost models for autoscheduling.
- [61] Samyam Rajbhandari, Jeff Rasley, Olatunji Rber, and Yuxiong He. 2020. ZERO: Memory Optimizations Toward Training Trillion Parameter Models. In *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis (SC)*. 1–16. <https://doi.org/10.1109/SC41405.2020.00024> DeepSpeed ZeRO optimizer partitioning for memory-efficient distributed training.
- [62] Mehdi Rakhshanfar and Aliakbar Zarandi. 2021. A Survey on Machine Learning-based Design Space Exploration for Processor Architectures. *Journal of Systems Architecture* 121 (2021), 102339. <https://doi.org/10.1016/j.sysarc.2021.102339>
- [63] Saeed Rashidi, Srinivas Srinivasan, Kazem Hamedani, and Tushar Krishna. 2020. ASTRA-SIM: Enabling SW/HW Co-Design Exploration for Distributed DL Training Platforms. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 81–92. <https://doi.org/10.1109/ISPASS48437.2020.00018>
- [64] Vijay Janapa Reddi et al. 2025. MLPerf Power: Benchmarking the Energy Efficiency of Machine Learning Inference. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 1–14. Energy efficiency benchmarking for ML inference workloads.
- [65] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmitt, Carole-Jean Wu, Brian Anderson, Maxim Breshekov, Mark Duber, et al. 2020. MLPerf Inference Benchmark. In *Proceedings of the 47th International Symposium on Computer Architecture (ISCA)*. 446–459. <https://doi.org/10.1109/ISCA45697.2020.00045> Standard ML inference benchmark suite with server and offline scenarios.
- [66] George F. Riley and Thomas R. Henderson. 2010. The ns-3 Network Simulator. *Modeling and Tools for Network Simulation* (2010), 15–34. [https://doi.org/10.1007/978-3-642-12331-3\\_2](https://doi.org/10.1007/978-3-642-12331-3_2)
- [67] Arun F. Rodrigues, K. Scott Hemmert, Brian W. Barrett, Chad Kersey, Ron Oldfield, Marlo Weston, R. Risen, Jeanine Cook, Paul Rosenfeld, Elliott Cooper-Balis, and Bruce Jacob. 2012. The Structural Simulation Toolkit. In *ACM SIGMETRICS Performance Evaluation Review*, Vol. 38. 37–42. <https://doi.org/10.1145/1964218.1964225> Modular framework for system-level simulation, widely used for HPC and interconnect modeling.
- [68] Ananda Samajdar, Yuhao Zhu, Paul Whatmough, Matthew Mattina, and Tushar Krishna. 2019. A Systematic Methodology for Characterizing Scalability of DNN Accelerators using SCALE-Sim. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 58–68. <https://doi.org/10.1109/ISPASS.2019.00016> Cycle-accurate systolic array simulator for DNN accelerator DSE.
- [69] Zhuomin Shen, Jaeho Kim, et al. 2025. AQUA: Network-Accelerated Memory Offloading for LLMs in Scale-Up GPU Domains. In *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 1–16. <https://doi.org/10.1145/3676641.3715983> Improves LLM inference responsiveness by 20x through network-accelerated memory offloading.
- [70] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. 2020. Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism. In *arXiv preprint arXiv:1909.08053*.

- 1045 Intra-layer tensor parallelism for large language model training.  
 1046 [71] Srinivas Sridharan, Taekyung Heo, Jinwoo Choi, Garyfallia Yu, Saeed Rashidi,  
 1047 William Won, Zhaodong Meng, and Tushar Krishna. 2023. Chakra: Advancing  
 1048 Performance Benchmarking and Co-design using Standardized Execution Traces.  
*arXiv preprint arXiv:2305.14516* (2023).
- 1049 [72] Foteini Strati, Zhendong Zhang, George Manos, Ixeia Sanchez Periz, Qinghai  
 1050 Hu, Tiancheng Chen, Berk Buzu, Song Han, Pamela Delgado, and Ana Klimovic.  
 1051 2025. Sailor: Automating Distributed Training over Dynamic, Heterogeneous,  
 1052 and Geo-distributed Clusters. In *Proceedings of the 30th ACM Symposium on  
 1053 Operating Systems Principles (SOSP)*. 1–18. Automated distributed training with  
 1054 runtime/memory simulation over heterogeneous resources. ETH Zurich/MIT.
- 1055 [73] Jonas Svedas, Hannah Watson, Nathan Laubeuf, Diksha Moolchandani, Abubakr  
 1056 Nada Arjun Singh, Dwaipayan Biswas, James Myers, and Debjyoti Bhattacharjee.  
 1057 2025. A Survey of End-to-End Modeling for Distributed DNN Training:  
 1058 Workloads, Simulators, and TCO. *arXiv preprint arXiv:2506.09275* (2025). Com-  
 1059 prehensive survey of distributed DNN training simulators covering workload  
 1060 representation, simulation infrastructure, and TCO/carbon modeling.
- 1061 [74] Ondrej Sykora, Alexis Rucker, Charith Mendis, Rajkishore Barik,  
 1062 Phitchaya Mangpo Phothilimthana, and Saman Amarasinghe. 2022. GRANITE:  
 1063 A Graph Neural Network Model for Basic Block Throughput Estimation. In  
 1064 *Proceedings of the IEEE International Symposium on Workload Characterization  
 1065 (IISWC)*. 1–13. <https://doi.org/10.1109/IISWC55918.2022.00014>
- 1066 [75] Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, and Joel S. Emer. 2017. Efficient  
 1067 Processing of Deep Neural Networks: A Tutorial and Survey. In *Proceedings  
 1068 of the IEEE*, Vol. 105. 2295–2329. <https://doi.org/10.1109/JPROC.2017.2761740>
- 1069 Canonical DNN accelerator taxonomy covering dataflows, data reuse, and energy  
 1070 efficiency.
- 1071 [76] Philippe Tillet, H. T. Kung, and David Cox. 2019. Triton: An Intermediate  
 1072 Language and Compiler for Tiled Neural Network Computations. In *Proceedings  
 1073 of the 3rd ACM SIGPLAN International Workshop on Machine Learning and  
 1074 Programming Languages (MAPL)*. 10–19. <https://doi.org/10.1145/3315508.3329973>
- 1075 Tile-based GPU programming with heuristic performance model for kernel  
 1076 generation.
- 1077 [77] Jan Treibig, Georg Hager, and Gerhard Wellein. 2010. LIKWID: A Lightweight  
 1078 Performance-Oriented Tool Suite for x86 Multicore Environments. In *Proceedings  
 1079 of the 39th International Conference on Parallel Processing Workshops (ICPPW)*. 207–  
 1080 216. <https://doi.org/10.1109/ICPPW.2010.38> Lightweight tools for thread/cache  
 1081 topology, affinity, and performance counter measurement.
- 1082 [78] Adrian Tschanz, Mohamed Awad, et al. 2025. SwizzlePerf: Hardware-Aware  
 1083 LLMs for GPU Kernel Performance Optimization. *arXiv preprint arXiv:2508.20258*  
 1084 (2025). LLM-based spatial optimization for GPU kernels, up to 2.06x speedup  
 1085 via swizzling.
- 1086 [79] Xizheng Wang, Qingxu Li, Yichi Xu, Gang Lu, Heyang Zhou, Sen Zhang, Yikai  
 1087 Zhu, Yang Liu, Pengcheng Zhang, Kun Qian, et al. 2025. SimAI: Unifying Ar-  
 1088 chitecture Design and Performance Tuning for Large-Scale LLM Training with  
 1089 Scalability and Precision. In *Proceedings of the 22nd USENIX Symposium on  
 1090 Networked Systems Design and Implementation (NSDI)*. 1–18. Full-stack LLM  
 1091 training simulator achieving 98.1% alignment with real-world results. Alibaba  
 1092 Cloud/Tsinghua..
- 1093 [80] Zixian Wang et al. 2025. SynPerf: Synthesizing High-Performance GPU Kernels  
 1094 via Pipeline Decomposition. *arXiv preprint* (2025). Under review.
- 1095 [81] Samuel Williams, Andrew Waterman, and David Patterson. 2009. Roofline: An  
 1096 Insightful Visual Performance Model for Multicore Architectures. *Commun.  
 1097 ACM* 52, 4 (2009), 65–76. <https://doi.org/10.1145/1498765.1498785>
- 1098 [82] William Won, Taekyung Heo, Saeed Rashidi, Saeed Talati, Srinivas Srinivasan,  
 1099 and Tushar Krishna. 2023. ASTRA-sim2.0: Modeling Hierarchical Networks and  
 1100 Disaggregated Systems for Large-Model Training at Scale. In *Proceedings of the  
 1101 IEEE International Symposium on Performance Analysis of Systems and Software  
 1102 (ISPASS)*. 283–294. <https://doi.org/10.1109/ISPASS57527.2023.00035>
- 1103 [83] Yannan Nellin Wu, Joel Emer, and Vivienne Sze. 2022. Sparseloop: An Analytical  
 1104 Approach to Sparse Tensor Accelerator Modeling. In *Proceedings of the 55th  
 1105 IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1–15. <https://doi.org/10.1109/MICRO56248.2022.000078>
- 1106 [84] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-  
 1107 Gon Chun. 2022. ORCA: A Distributed Serving System for Transformer-Based  
 1108 Generative Models. In *Proceedings of the 16th USENIX Symposium on Operating  
 1109 Systems Design and Implementation (OSDI)*. 521–538.
- 1110 [85] Geoffrey X. Yu, Yubo Gao, Pavel Golber, and Asaf Cidon. 2021. Habitat: A  
 1111 Runtime-Based Computational Performance Predictor for Deep Neural Network  
 1112 Training. In *Proceedings of the USENIX Annual Technical Conference (ATC)*. 503–  
 1113 521.
- 1114 [86] Yi Zhai, Yu Cheng Wang, Peng Jiang, and Congming Kang. 2023. TLP: A Deep  
 1115 Learning-based Cost Model for Tensor Program Tuning. In *Proceedings of the  
 1116 28th ACM International Conference on Architectural Support for Programming  
 1117 Languages and Operating Systems (ASPLOS)*. 833–845. <https://doi.org/10.1145/3575693.3575736>

1103  
 1104  
 1105  
 1106  
 1107  
 1108  
 1109  
 1110  
 1111  
 1112  
 1113  
 1114  
 1115  
 1116  
 1117  
 1118  
 1119  
 1120  
 1121  
 1122  
 1123  
 1124  
 1125  
 1126  
 1127  
 1128  
 1129  
 1130  
 1131  
 1132  
 1133  
 1134  
 1135  
 1136  
 1137  
 1138  
 1139  
 1140  
 1141  
 1142  
 1143  
 1144  
 1145  
 1146  
 1147  
 1148  
 1149  
 1150  
 1151  
 1152  
 1153  
 1154  
 1155  
 1156  
 1157  
 1158  
 1159  
 1160